

Universidad Nacional de San Agustín

# Base de datos

## Spanner: Becoming a SQL System

MSc. Vicente Machaca Arceda

2021



## Spanner

What is spanner?

Architecture

Scalability and Disponibility

# Spanner

What is spanner?



“Spanner is Google’s **scalable, multiversion, globally distributed, and synchronously replicated database**” [1].

Spanner is used as an OLTP database management system ( for AdWords and Google Play), and is publicly available in Cloud Spanner on the Google Cloud Platform (GCP) [2].



Cloud  
Spanner

# Spanner

What is spanner?



	Spanner	Traditional relational	Traditional non-relational
Scheme	yes	yes	no
SQL	yes	yes	no
Consistency	strong	strong	eventual
Availavility	high	Failover	High
Scalability	Horizontal	Vertical	Horizontal
Replication	Automatic	Configurable	Configurable

Table: Properties of Spanner. Source: Cloud spanner


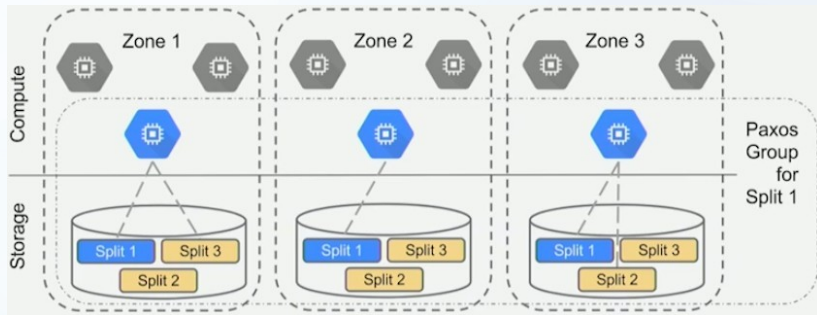


Figure: Example of horizontal partition.

# Architecture

Sharded, geo-replicated relational database



**Figure:** Spanner is a sharded, geo-replicated relational database. It uses a replicated write-ahead redo log, and the Paxos protocol.

*Consistent read*

```
SELECT * FROM Event WHERE Name = 'Devovx PL';
```

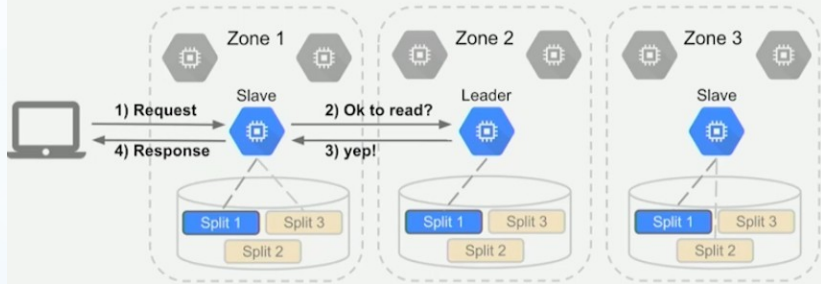


Figure: Slave is updated.

# Architecture

## Consistent read



*Consistent read*

```
SELECT * FROM Event WHERE Name = 'Devoxx PL';
```

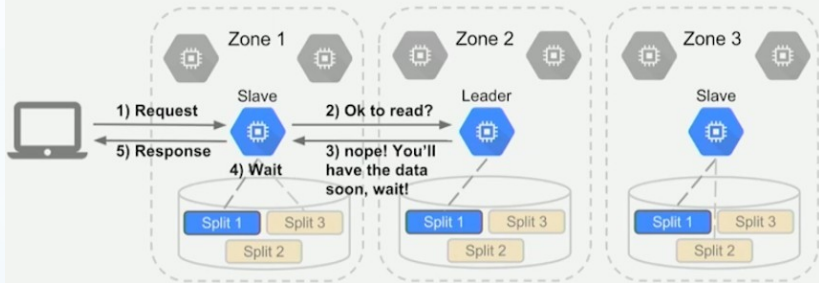


Figure: Slave is not updated.



### Read-Write Transaction

```
m = append(m, spanner.InsertOrUpdate("Talks", []string{"Name", "Time"},  
                                     []interface{}{"Cloud Spanner 101", "2017-06-22T13:20+2"}))  
txn.BufferWrite(m)
```

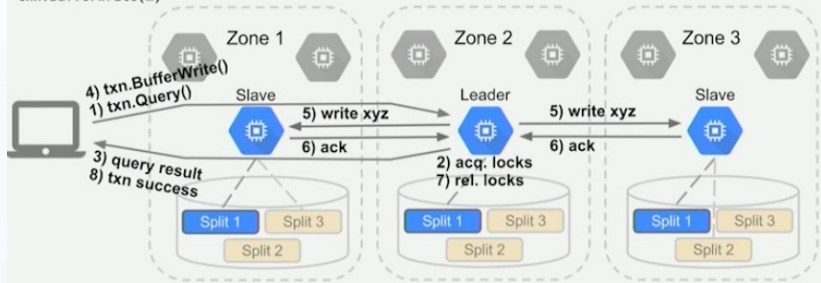


Figure: Read-write transaction.



## Why Spanner is scalable?

Spanner follows a range sharding architecture, so it reduce disk space and memory use.

## What technologies support Spanner's disponibility?

Truetime [3], automatic replication.



- [1] J. C. Corbett, J. Dean, M. Epstein, A. Fikes, C. Frost, J. J. Furman, S. Ghemawat, A. Gubarev, C. Heiser, P. Hochschild *et al.*, “Spanner: Google’s globally distributed database,” *ACM Transactions on Computer Systems (TOCS)*, vol. 31, no. 3, pp. 1–22, 2013.
- [2] D. F. Bacon, N. Bales, N. Bruno, B. F. Cooper, A. Dickinson, A. Fikes, C. Fraser, A. Gubarev, M. Joshi, E. Kogan *et al.*, “Spanner: Becoming a sql system,” in *Proceedings of the 2017 ACM International Conference on Management of Data*, 2017, pp. 331–343.
- [3] E. Brewer, “Spanner, truetime and the cap theorem,” 2017.

